

# q1-svm

November 6, 2023

## 0.1 Assignment 3

## 0.2 Support Vector Machine

### 0.2.1 Name: Gautam Kumar

### 0.2.2 Roll Number: 21CS30020

```
[24]: from ucimlrepo import fetch_ucirepo
import pandas as pd
```

```
spambase = fetch_ucirepo(id=94)
# data (as pandas dataframes)
X = spambase.data.features
y = spambase.data.targets
# metadata
print(spambase.metadata)
# variable information
print(spambase.variables)
# loading as dataframe
x = spambase.data.features
y = spambase.data.targets
```

```
{'uci_id': 94, 'name': 'Spambase', 'repository_url':
'https://archive.ics.uci.edu/dataset/94/spambase', 'data_url':
'https://archive.ics.uci.edu/static/public/94/data.csv', 'abstract':
'Classifying Email as Spam or Non-Spam', 'area': 'Computer Science', 'tasks':
['Classification'], 'characteristics': ['Multivariate'], 'num_instances': 4601,
'num_features': 57, 'feature_types': ['Integer', 'Real'], 'demographics': [],
'target_col': ['Class'], 'index_col': None, 'has_missing_values': 'no',
'missing_values_symbol': None, 'year_of_dataset_creation': 1999, 'last_updated':
'Mon Aug 28 2023', 'dataset_doi': '10.24432/C53G6X', 'creators': ['Mark
Hopkins', 'Erik Reeber', 'George Forman', 'Jaap Suermondt'], 'intro_paper':
None, 'additional_info': {'summary': 'The "spam" concept is diverse:
advertisements for products/web sites, make money fast schemes, chain letters,
pornography...\n\nThe classification task for this dataset is to determine
whether a given email is spam or not.\n\t\nOur collection of spam e-mails came
from our postmaster and individuals who had filed spam. Our collection of non-
spam e-mails came from filed work and personal e-mails, and hence the word
\'george\' and the area code \'650\' are indicators of non-spam. These are
```

useful when constructing a personalized spam filter. One would either have to blind such non-spam indicators or get a very wide collection of non-spam to generate a general purpose spam filter.

For background on spam: Cranor, Lorrie F., LaMacchia, Brian A. Spam!, Communications of the ACM, 41(8):74-83, 1998.

Typical performance is around ~7% misclassification error. False positives (marking good mail as spam) are very undesirable. If we insist on zero false positives in the training/testing set, 20-25% of the spam passed through the filter. See also Hewlett-Packard Internal-only Technical Report. External version forthcoming.

```
{
  'purpose': None, 'funded_by': None,
  'instances_represent': 'Emails', 'recommended_data_splits': None,
  'sensitive_data': None, 'preprocessing_description': None, 'variable_info': 'The
last column of \'spambase.data\' denotes whether the e-mail was considered spam
(1) or not (0), i.e. unsolicited commercial e-mail. Most of the attributes
indicate whether a particular word or character was frequently occurring in the
e-mail. The run-length attributes (55-57) measure the length of sequences of
consecutive capital letters. For the statistical measures of each attribute,
see the end of this file. Here are the definitions of the attributes:\r\n\r\n48
continuous real [0,100] attributes of type word_freq_WORD \r\n= percentage of
words in the e-mail that match WORD, i.e. 100 * (number of times the WORD
appears in the e-mail) / total number of words in e-mail. A "word" in this case
is any string of alphanumeric characters bounded by non-alphanumeric characters
or end-of-string.\r\n\r\n6 continuous real [0,100] attributes of type
char_freq_CHAR] \r\n= percentage of characters in the e-mail that match CHAR,
i.e. 100 * (number of CHAR occurrences) / total characters in e-mail\r\n\r\n1
continuous real [1,...] attribute of type capital_run_length_average \r\n=
average length of uninterrupted sequences of capital letters\r\n\r\n1 continuous
integer [1,...] attribute of type capital_run_length_longest \r\n= length of
longest uninterrupted sequence of capital letters\r\n\r\n1 continuous integer
[1,...] attribute of type capital_run_length_total \r\n= sum of length of
uninterrupted sequences of capital letters \r\n= total number of capital letters
in the e-mail\r\n\r\n1 nominal {0,1} class attribute of type spam\r\n= denotes
whether the e-mail was considered spam (1) or not (0), i.e. unsolicited
commercial e-mail. \r\n', 'citation': None}}
```

	name	role	type	demographic	\
0	word_freq_make	Feature	Continuous	None	
1	word_freq_address	Feature	Continuous	None	
2	word_freq_all	Feature	Continuous	None	
3	word_freq_3d	Feature	Continuous	None	
4	word_freq_our	Feature	Continuous	None	
5	word_freq_over	Feature	Continuous	None	
6	word_freq_remove	Feature	Continuous	None	
7	word_freq_internet	Feature	Continuous	None	
8	word_freq_order	Feature	Continuous	None	
9	word_freq_mail	Feature	Continuous	None	
10	word_freq_receive	Feature	Continuous	None	
11	word_freq_will	Feature	Continuous	None	
12	word_freq_people	Feature	Continuous	None	
13	word_freq_report	Feature	Continuous	None	

14	word_freq_addresses	Feature	Continuous	None
15	word_freq_free	Feature	Continuous	None
16	word_freq_business	Feature	Continuous	None
17	word_freq_email	Feature	Continuous	None
18	word_freq_you	Feature	Continuous	None
19	word_freq_credit	Feature	Continuous	None
20	word_freq_your	Feature	Continuous	None
21	word_freq_font	Feature	Continuous	None
22	word_freq_000	Feature	Continuous	None
23	word_freq_money	Feature	Continuous	None
24	word_freq_hp	Feature	Continuous	None
25	word_freq_hpl	Feature	Continuous	None
26	word_freq_george	Feature	Continuous	None
27	word_freq_650	Feature	Continuous	None
28	word_freq_lab	Feature	Continuous	None
29	word_freq_labs	Feature	Continuous	None
30	word_freq_telnet	Feature	Continuous	None
31	word_freq_857	Feature	Continuous	None
32	word_freq_data	Feature	Continuous	None
33	word_freq_415	Feature	Continuous	None
34	word_freq_85	Feature	Continuous	None
35	word_freq_technology	Feature	Continuous	None
36	word_freq_1999	Feature	Continuous	None
37	word_freq_parts	Feature	Continuous	None
38	word_freq_pm	Feature	Continuous	None
39	word_freq_direct	Feature	Continuous	None
40	word_freq_cs	Feature	Continuous	None
41	word_freq_meeting	Feature	Continuous	None
42	word_freq_original	Feature	Continuous	None
43	word_freq_project	Feature	Continuous	None
44	word_freq_re	Feature	Continuous	None
45	word_freq_edu	Feature	Continuous	None
46	word_freq_table	Feature	Continuous	None
47	word_freq_conference	Feature	Continuous	None
48	char_freq_;	Feature	Continuous	None
49	char_freq(	Feature	Continuous	None
50	char_freq[	Feature	Continuous	None
51	char_freq!	Feature	Continuous	None
52	char_freq\$	Feature	Continuous	None
53	char_freq#	Feature	Continuous	None
54	capital_run_length_average	Feature	Continuous	None
55	capital_run_length_longest	Feature	Continuous	None
56	capital_run_length_total	Feature	Continuous	None
57	Class	Target	Binary	None

	description	units	missing_values
0	None	None	no
1	None	None	no

2	None	None	no
3	None	None	no
4	None	None	no
5	None	None	no
6	None	None	no
7	None	None	no
8	None	None	no
9	None	None	no
10	None	None	no
11	None	None	no
12	None	None	no
13	None	None	no
14	None	None	no
15	None	None	no
16	None	None	no
17	None	None	no
18	None	None	no
19	None	None	no
20	None	None	no
21	None	None	no
22	None	None	no
23	None	None	no
24	None	None	no
25	None	None	no
26	None	None	no
27	None	None	no
28	None	None	no
29	None	None	no
30	None	None	no
31	None	None	no
32	None	None	no
33	None	None	no
34	None	None	no
35	None	None	no
36	None	None	no
37	None	None	no
38	None	None	no
39	None	None	no
40	None	None	no
41	None	None	no
42	None	None	no
43	None	None	no
44	None	None	no
45	None	None	no
46	None	None	no
47	None	None	no
48	None	None	no
49	None	None	no

50		None	None	no
51		None	None	no
52		None	None	no
53		None	None	no
54		None	None	no
55		None	None	no
56		None	None	no
57	spam (1) or not spam (0)	None		no

```
[25]: import numpy as np
      from sklearn.model_selection import train_test_split
      import matplotlib.pyplot as plt
      from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
[26]: from sklearn import datasets
      from sklearn.model_selection import train_test_split
      from sklearn.svm import SVC
```

```
[27]: print(len(X))
```

4601

```
[28]: y = y.iloc[:, 0].ravel()
      X_train, X_test, y_train, y_test = train_test_split(X, y,
      test_size=0.2, random_state=42)
```

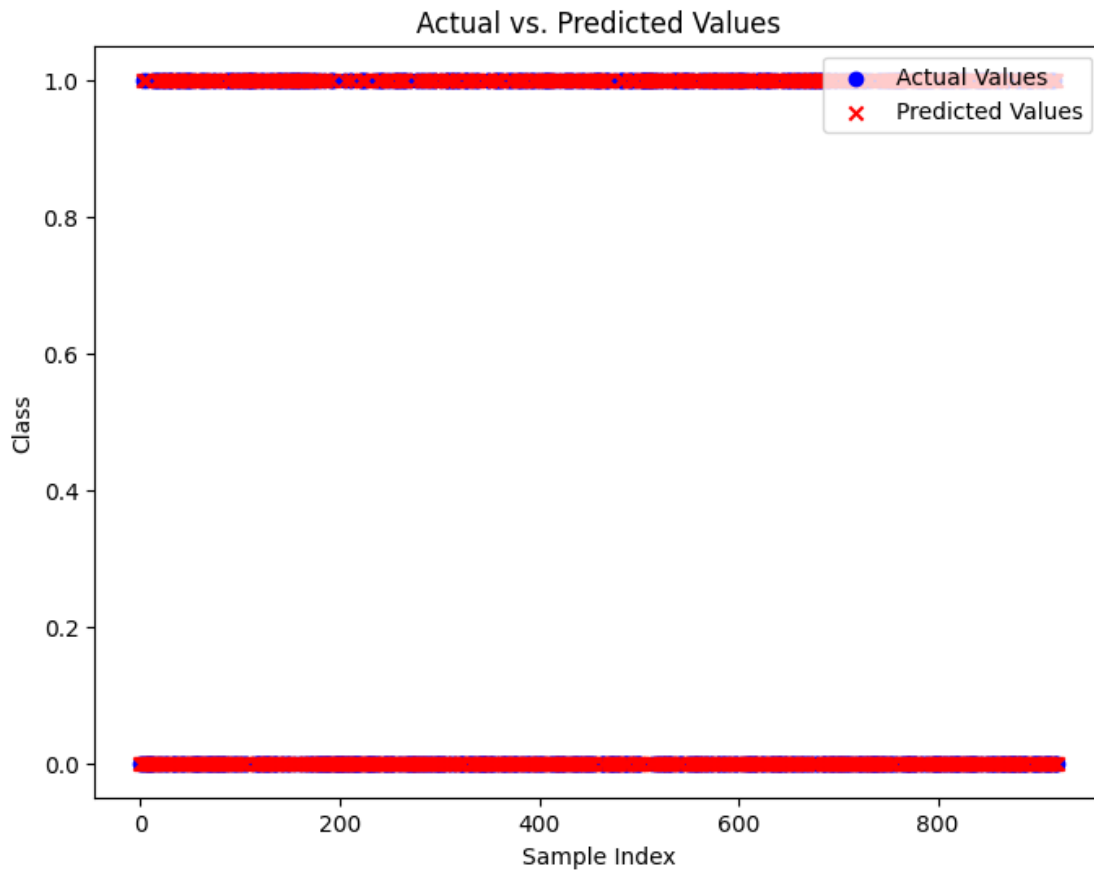
```
[29]: svm_model = SVC(kernel='linear')
      svm_model.fit(X_train, y_train)
```

```
[29]: SVC(kernel='linear')
```

```
[30]: y_pred = svm_model.predict(X_test)
```

```
[31]: ##Accuracy without applying Regularisation
      accuracy = accuracy_score(y_test, y_pred)
```

```
[32]: plt.figure(figsize=(8, 6))
      plt.scatter(range(len(y_test)), y_test, color='blue', label='Actual Values',
      marker='o')
      plt.scatter(range(len(y_test)), y_pred, color='red', label='Predicted Values',
      marker='x')
      plt.title('Actual vs. Predicted Values')
      plt.xlabel('Sample Index')
      plt.ylabel('Class')
      plt.legend(loc='upper right')
      plt.show()
```



```
[33]: C_values = [0.001, 0.1, 1, 10, 100]
accuracy_scores = []
y_preds = []
```

```
[34]: for C in C_values:
    svm_model = SVC(kernel='linear', C=C)
    svm_model.fit(X_train, y_train)
    y_pred = svm_model.predict(X_test)
    y_preds.append(y_pred)
    accuracy = accuracy_score(y_test, y_pred)
    accuracy_scores.append(accuracy)
```

```
[35]: accuracy_table = pd.DataFrame({'C': C_values, 'Accuracy': accuracy_scores})
```

```
[36]: accuracy_table
```

```
[36]:      C  Accuracy
0  0.001  0.865364
1  0.100  0.925081
```

```


2    1.000  0.922910
3   10.000  0.918567
4  100.000  0.913138


```

```


[37]: i = 0
      for y_pred in y_preds:
          plt.figure(figsize=(8, 6))
          plt.scatter(range(len(y_test)), y_test, color='blue', label='Actual_
↪Values', marker='o')
          plt.scatter(range(len(y_test)), y_pred, color='red', label='Predicted_
↪Values', marker='x')
          plt.title(f'Actual vs. Predicted Values (Accuracy: {accuracy_scores[i]})')
          plt.xlabel('Sample Index')
          plt.ylabel('Class')
          plt.legend(loc='upper right')
          plt.show()
          i += 1


```






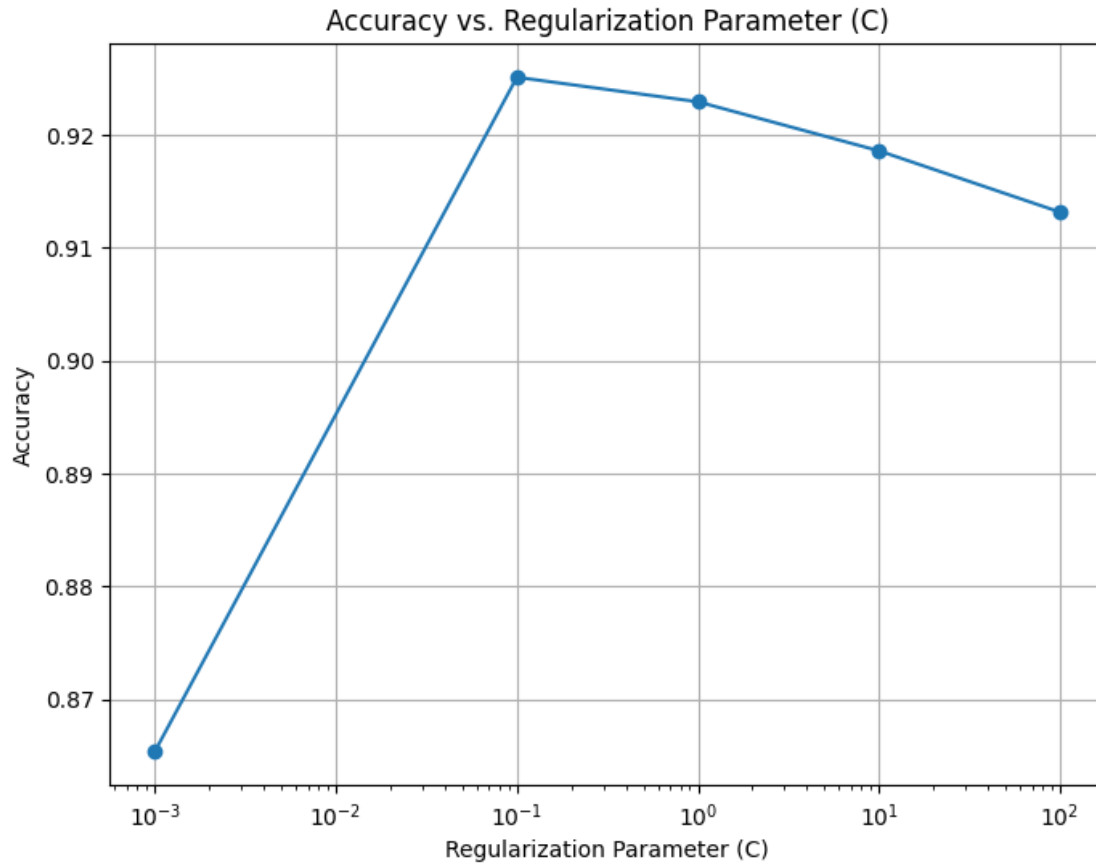








```
[38]: plt.figure(figsize=(8, 6))
plt.plot(C_values, accuracy_scores, marker='o')
plt.title('Accuracy vs. Regularization Parameter (C)')
plt.xlabel('Regularization Parameter (C)')
plt.xscale('log')
plt.ylabel('Accuracy')
plt.grid(True)
plt.show()
```



```
[39]: # Define a list of kernels to test
kernels = ['poly', 'poly', 'sigmoid', 'rbf']
degrees = [2, 3, 0, 0] # Specify the degrees for the polynomial kernels
```

```
[40]: results = []
for kernel, degree in zip(kernels, degrees):
    if kernel == 'poly':
        # Polynomial kernel
        svm_model = SVC(kernel=kernel, degree=degree)
    else:
        # Sigmoid or RBF kernel
        svm_model = SVC(kernel=kernel)

    svm_model.fit(X_train, y_train)
    y_pred = svm_model.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred, average='weighted',
    ↪ zero_division=0)
```

```

recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)
f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)

results.append({
    'Kernel': kernel,
    'Degree': degree if kernel == 'poly' else None,
    'Accuracy': accuracy,
    'Precision': precision,
    'Recall': recall,
    'F1 Score': f1
})
results_df = pd.DataFrame(results)

```

```
[41]: results_df
```

```
[41]:
```

	Kernel	Degree	Accuracy	Precision	Recall	F1 Score
0	poly	2.0	0.649294	0.731458	0.649294	0.577828
1	poly	3.0	0.625407	0.718539	0.625407	0.533076
2	sigmoid	NaN	0.635179	0.632524	0.635179	0.633420
3	rbf	NaN	0.662324	0.662088	0.662324	0.644053

```
[42]: degrees = [1, 1, 3, 3]
C_values = [0.01, 100, 0.01, 100]
train_accuracies = []
test_accuracies = []

```

```
[43]: for degree, C in zip(degrees, C_values):
    svm_model = SVC(kernel='poly', degree=degree, C=C)
    svm_model.fit(X_train, y_train)

    # Training accuracy
    y_train_pred = svm_model.predict(X_train)
    train_accuracy = accuracy_score(y_train, y_train_pred)
    train_accuracies.append(train_accuracy)

    # Test accuracy
    y_test_pred = svm_model.predict(X_test)
    test_accuracy = accuracy_score(y_test, y_test_pred)
    test_accuracies.append(test_accuracy)

```

```
[44]: experiment_results = pd.DataFrame({
    'Polynomial Degree': degrees,
    'Regularization Parameter (C)': C_values,
    'Train Accuracy': train_accuracies,
    'Test Accuracy': test_accuracies
})

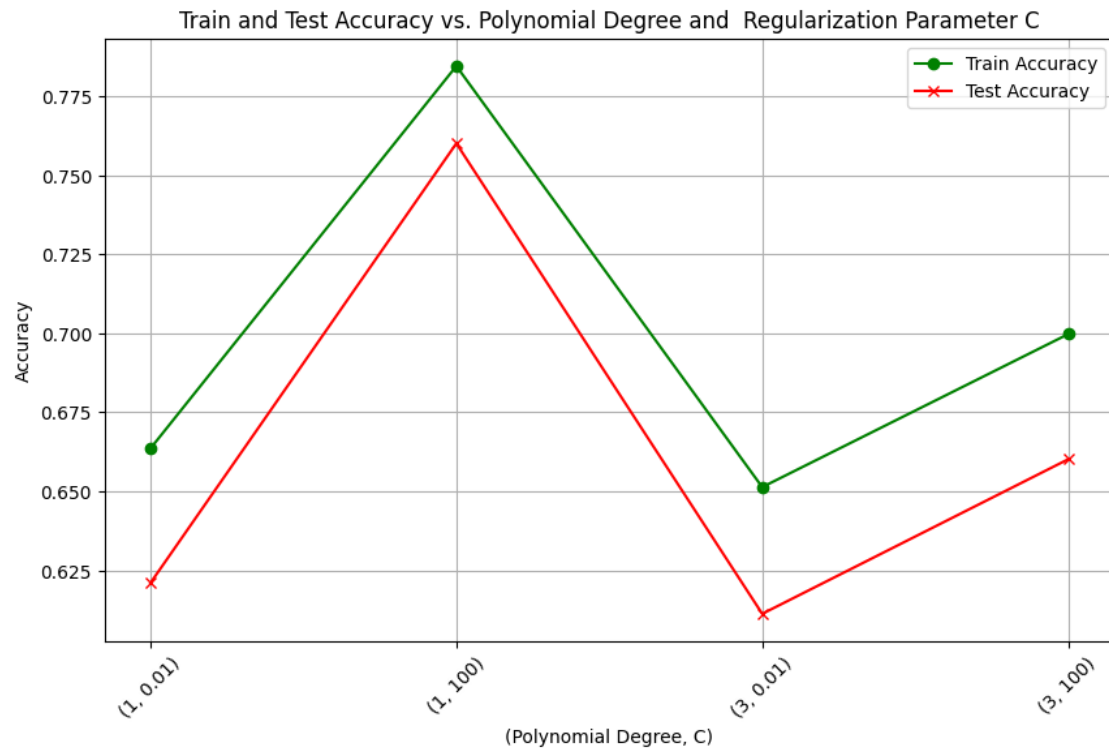
```

```
[45]: experiment_results
```

```
[45]: Polynomial Degree Regularization Parameter (C) Train Accuracy \
0      1      0.01      0.663587
1      1      100.00     0.784511
2      3      0.01      0.651359
3      3      100.00     0.699728

Test Accuracy
0      0.621064
1      0.760043
2      0.611292
3      0.660152
```

```
[46]: plt.figure(figsize=(10, 6))
plt.plot(range(len(experiment_results)), train_accuracies,color='green',
        ↪marker='o', label='Train Accuracy')
plt.plot(range(len(experiment_results)), test_accuracies,color='red',
        ↪marker='x', label='Test Accuracy')
plt.xticks(range(len(experiment_results)), [f'({degree}, {C})' for degree, C in
        ↪zip(degrees, C_values)], rotation=45)
plt.xlabel('(Polynomial Degree, C)')
plt.ylabel('Accuracy')
plt.title('Train and Test Accuracy vs. Polynomial Degree and Regularization
        ↪Parameter C')
plt.legend()
plt.grid(True)
plt.show()
```



[ ]:

[ ]: